Venkata Sai Sravya Sambaturu

sravya.sambaturu75@gmail.com | github.com/sravyasambaturu | linkedin.com/in/sravyasa | Farmington Hills, Michigan

PROFESSIONAL SUMMARY

DevOps / Cloud / Site Reliability Engineer with 5+ years of experience designing and automating scalable, multi-cloud (AWS, GCP, Azure) infrastructure. Skilled in Terraform, Kubernetes, CI/CD pipeline engineering, monitoring/observability, and workflow automation, with a track record of reducing deployment time, improving reliability, and enabling developer self-service at scale.

EDUCATION

M.S. Computer Science & Information Systems - University of Michigan-Flint Jan 2025 - Dec 2026
GPA: 4.0/4.0

B.Tech. Mechanical Engineering - Vignan Institute of Technology and Sciences Aug 2015 - May 2019

TECHNICAL SKILLS

Programming & Cloud: Python, Java, Shell/Bash, PowerShell, YAML, Groovy, AWS, GCP, Azure

Infrastructure & Automation: Terraform, Ansible, Docker, Kubernetes, OpenShift, Helm, VPC, Load Balancers, Service Mesh

CI/CD & DevOps: Git, Jenkins, GitHub Actions, GitLab CI/CD, ArgoCD (GitOps), Nginx, Apache

Monitoring & Security: Prometheus, Grafana, ELK Stack, Splunk, Datadog, SonarQube, SAST/DAST

Databases & Storage: SQL, MongoDB, MySQL, Oracle DB, AWS RDS, S3

Certifications: GCP Cloud Engineer & Architect, AWS Cloud Practitioner, HashiCorp Terraform Associate, GCP ML Engineer

PROFESSIONAL EXPERIENCE

Solutions Engineering Intern | ThyssenKrupp Materials **Jul 2026 - Present**

- Gaining hands-on experience with Microsoft Azure services, SQL database development, and workflow automation using Azure Functions and Power Automate to support business process improvements
- Collaborating with cross-functional teams to design, test, and document automated workflows, applying SQL querying and data handling in a live enterprise environment

Graduate Research Assistant | University of Michigan **May 2025 - June 2026**

- Building an MLOps pipeline (MLflow) with containerized, automated training and CI/CD-style deployment, reducing iteration time from 2.5 hours to 10 minutes

DevOps Engineer | London Stock Exchange Group **Dec 2023 - Jul 2024**

- Automated infrastructure provisioning and API deployments on AWS for 15+ microservices using Terraform, Python, and GitLab CI/CD, cutting deployment time from 4 hours to 45 minutes and eliminating 80% of manual errors
- Led migration of 25+ CI/CD pipelines from Jenkins to GitLab (Java, .NET, Node.js) using Groovy/YAML and AWS-integrated GitLab Runners, reducing pipeline failures by 60%
- Delivered GitLab CI/CD training to 40+ engineers across 5 teams, cutting onboarding time from 2 weeks to 3 days and increasing deployment frequency 3x in one quarter

GCP DevOps Engineer | Techolution **Jan 2022 - Dec 2023**

- Deployed GCP infrastructure for 4 enterprise clients (Wendy's, Main Street Bank, HarperCollins, Artajasa) using Terraform, GitLab CI/CD, and Kubernetes, cutting provisioning time from 2 weeks to 2 days for 50+ self-service developers
- Containerized Java/Python workloads (60% smaller images) and built GitHub Actions mobile pipelines and automated Vertex AI deployments, reducing deployment time from 3 days to 4 hours
- Ran a DevOps enablement program (8 workshops on GitLab, Docker, Kubernetes, OpenShift) for 100+ engineers, cutting new-hire setup time 70% and support tickets 40%

Software Associate - Alexa | Amazon Development Centre **Oct 2019 - Dec 2021**

- Supported AWS-based ML infrastructure for Alexa, managing 10,000+ audio/transcribed files in S3 with 99.5% uptime and GDPR compliance
- Optimized data workflows as SME for 15+ team members, reducing processing errors 35% and cutting incident resolution time from 4 hours to 1 hour and receiving the Extra Mile Award

TECHNICAL PROJECTS

Multi-Cloud LLM Serving Platform: Deployed Llama 2 7B across AWS/GCP/Azure with Kubernetes and vLLM (2000+ tokens/sec); built Prometheus/Grafana cost monitoring, cutting multi-cloud inference costs 40% via spot instances

Full-Stack E-Commerce, Multi-AZ Infrastructure: React + Spring Boot microservices deployed on AWS EKS across 2 AZs with Terraform/Kubernetes auto-scaling, Multi-AZ RDS, and Redis caching (150ms latency)